

# Real-time Rendering of Translucent Materials with Directional Subsurface Scattering

Alessandro Dal Corso

DTU Compute

7<sup>th</sup> November 2014



# Introduction



# Problem statement

Our goal is to represent SS phenomena in a synthetically generated image:

- Using a analytical BSSRDF model [Frisvad et al., 2014]
- Visually close as possible to a offline-rendered solution
- That does not need uv-mapping (only position data)
- In real-time or at least at interactive frame rates using a GPUs

# The rendering equation

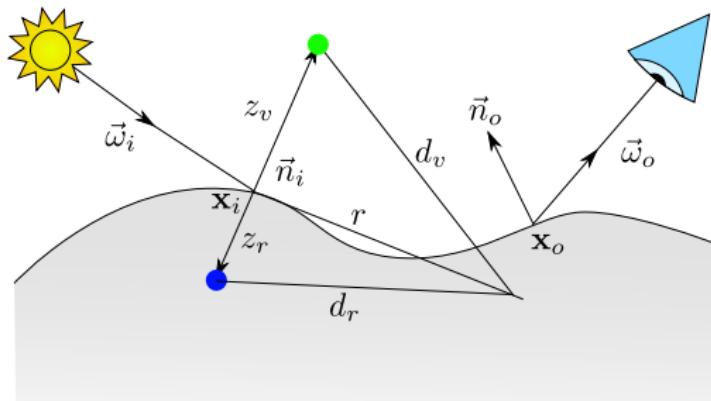
- We use the *area formulation* of the *rendering equation* [Jensen et al., 2001]:

$$L_o(\mathbf{x}_o, \vec{\omega}_o) = L_e(\mathbf{x}_o, \vec{\omega}_o) + \int_A \int_{2\pi} S(\mathbf{x}_i, \vec{\omega}_i, \mathbf{x}_o, \vec{\omega}_o) L_i(\mathbf{x}_i, \vec{\omega}_i) (\vec{n} \cdot \vec{\omega}_i) d\vec{\omega}_i dA_i$$

- Many BSSRDF functions have been proposed in literature [Jensen et al., 2001; D'Eon and Irving, 2011; Frisvad et al., 2014]

# Standard dipole

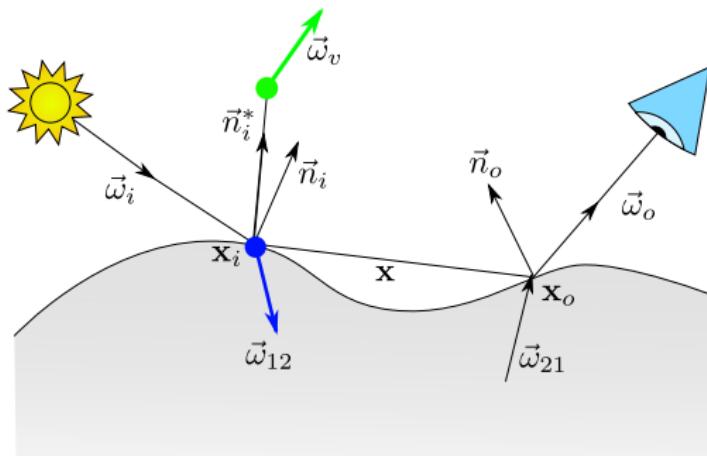
- The first BSSRDF model to be proposed for subsurface scattering (depends only on  $r$ )
- It models the interaction as two small light point sources, a **virtual** and a **real** source



$$S_d(\mathbf{x}_i, \vec{\omega}_i, \mathbf{x}_r, \vec{\omega}_o) = \frac{\alpha'}{4\pi^2} \left[ \frac{z_r(1 + \sigma_{tr}d_r) e^{-\sigma_{tr}d_r}}{d_r^3} + \frac{z_v(1 + \sigma_{tr}d_v) e^{-\sigma_{tr}d_v}}{d_v^3} \right]$$

# Directional dipole

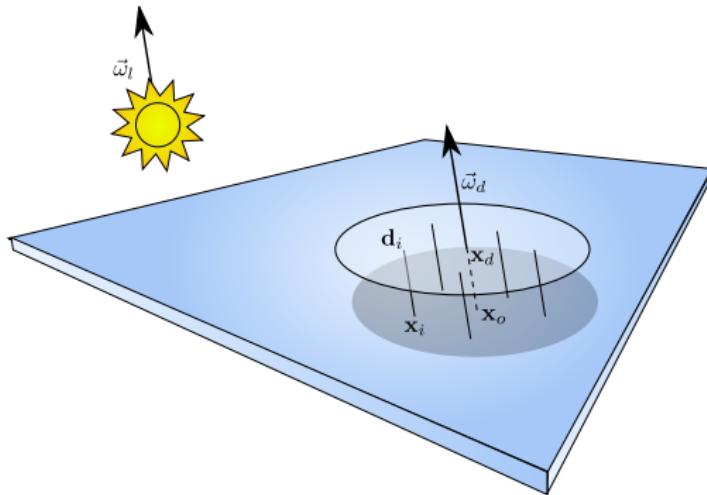
- [Frisvad et al., 2014] defined a new BSSRDF function that keeps the directionality of the incoming light into account
- Depends on the direction of the incoming light  $\vec{\omega}_i$ , the two points  $\mathbf{x}_o$  and  $\mathbf{x}_i$  and the two normals



$$S_d(\mathbf{x}_i, \vec{\omega}_i, \mathbf{x}_o) = S'_d(\mathbf{x}_o - \mathbf{x}_i, \vec{\omega}_{12}, d_r) - S'_d(\mathbf{x}_o - \mathbf{x}_v, \vec{\omega}_v, d_v)$$

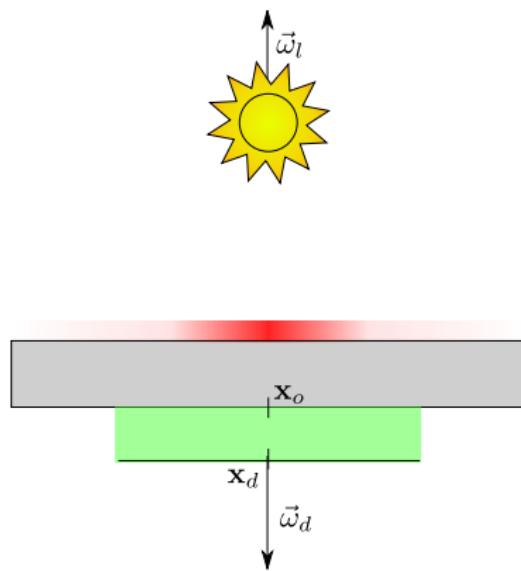
# Method

- Scattering effects have a limited range, that depends on the scattering properties of the material (especially on  $1/\sigma_{tr}$ )
- We can then consider contributions from a disk on the surface
- The disk has a center point  $\mathbf{x}_d$  and a direction  $\vec{\omega}_d$

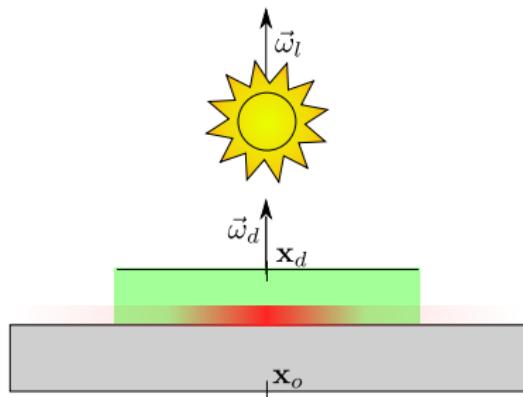


# Method

- We place the disk oriented towards the light ( $\vec{\omega}_d = \vec{\omega}_l$ ) and close enough to the light in order to cover the surface



*Naïve disk placement*



*Our disk placement*

# Method

- Given these assumptions on the disk, we discretize the rendering equation (using uniform sampling):

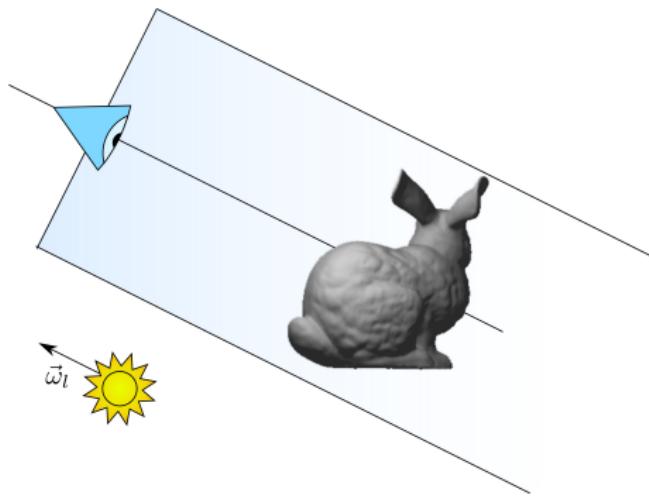
$$L_o(\mathbf{x}_o, \vec{\omega}_o) = L_e(\mathbf{x}_o, \vec{\omega}_o) + \int_A \int_{2\pi} S(\mathbf{x}_i, \vec{\omega}_i, \mathbf{x}_o, \vec{\omega}_o) L_i(\mathbf{x}_i, \vec{\omega}_i) (\vec{n} \cdot \vec{\omega}_i) d\vec{\omega}_i dA_i$$

$$L_o(\mathbf{x}_o, \vec{\omega}_o) = L_d \frac{A_c}{N} \sum_{i=1}^N S(\mathbf{x}_i, \vec{\omega}_d, \mathbf{x}_o, \vec{\omega}_o) (\vec{\omega}_d \cdot \vec{n}_i)$$

- A uniform sampling is not enough, so we need to employ a particular sampling scheme

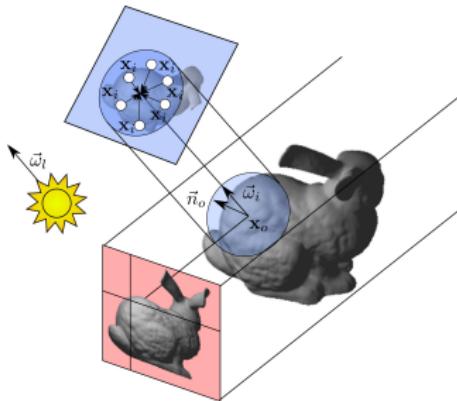
# Implementation (step 1)

- We render the scene from the light point of view (as in *shadow mapping*)
- We store positions and normals in a texture, the *light map*
- We get the closest points to light



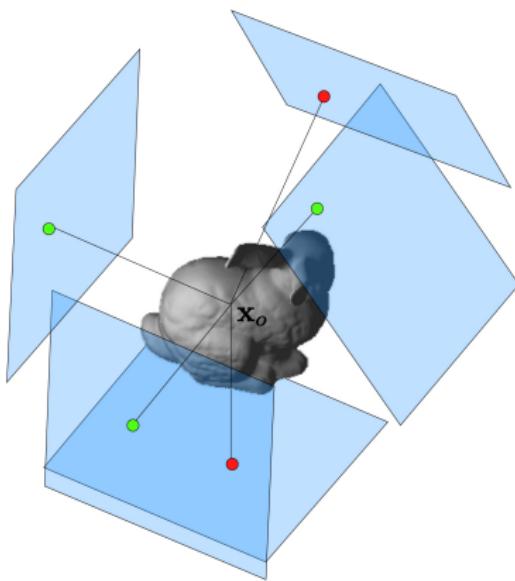
# Implementation (step 2)

- We render the object from a certain number of directions in the *radiance map*
- The directions are chosen randomly
- For each pixel we sample the points from the light map and sum up the BSSRDF contribution
- The result is accumulated over multiple frames to reduce noise

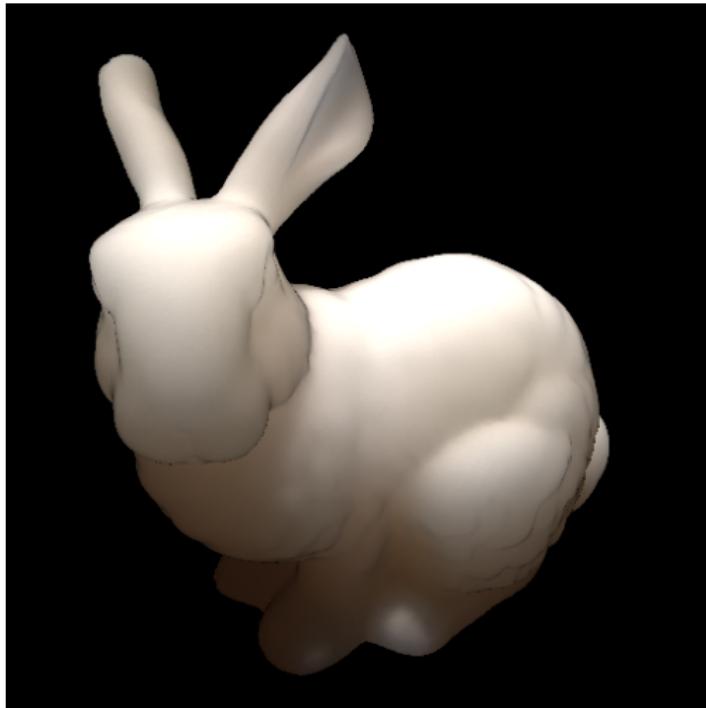


# Implementation (step 3)

- We finally sample the radiance map to get the single contribution for a point on the surface
- The result is averaged over the directions from which the surface point is visible



# Implementation



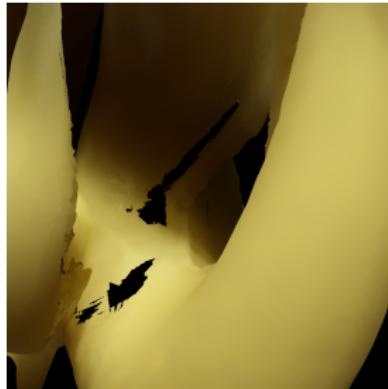
**Figure:** Stanford Bunny, potato material. Note the self shadowing automatically generated by the algorithm.

# Advantages

- Accounts for self occlusion
- Accounts for occlusion between different objects
- Directly coupled with an existing shadow mapping pipeline
- Low memory requirements compared to a full voxelization
- No UV mapping is necessary
- The final step can be adapted to forward and deferred shading pipelines

# Disadvantages

- Noise, partially overcome by sampling
- Cameras that cover the surface need to be placed manually (to avoid tearing)
- Inherited problems from shadow mapping (constant shadow bias)



# Results (quality)



(a) 1 frame (5 samples)

(b) ca. 100 frames

**Figure:** Result comparison for the Stanford Bunny model at different times.  
Parameters for potato.

# Extensions

- Multiple lights

We sum the contribution of multiple lights in the shader

- Point lights

We scale the incoming intensity by a inverse square distance term

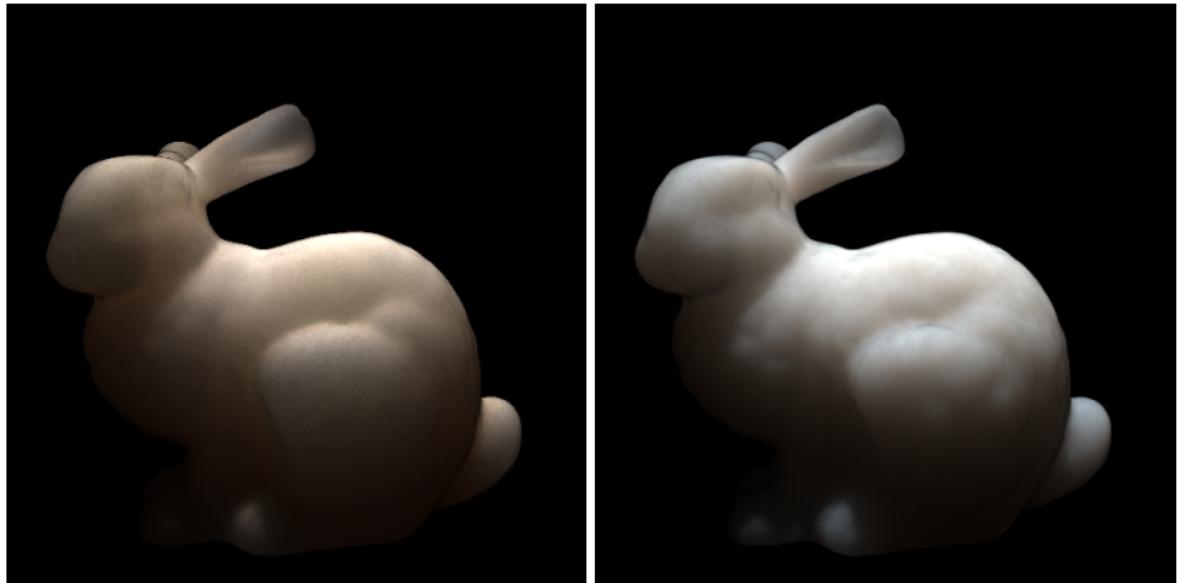


# Environment lighting

- We simulate it using 16 directional lights
- We choose "random" points on a environment map
- The distribution chosen accounts to make the points fall in areas with high luminosity



# Results (quality)

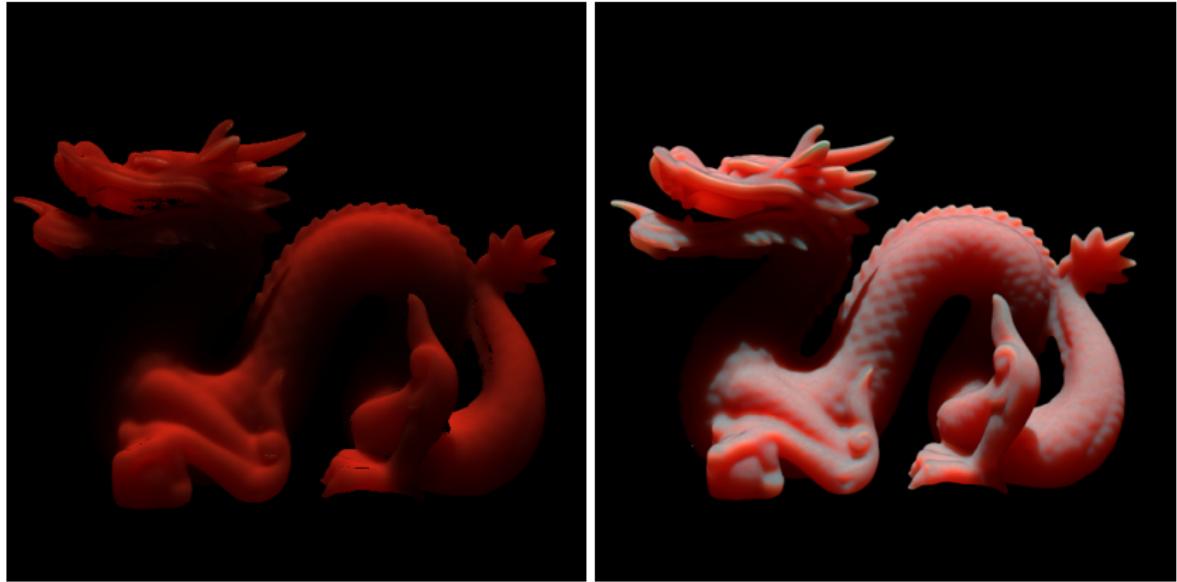


(a)Our method

(b)Reference

**Figure:** Result comparison for the Stanford Bunny model. Parameters for marble.

# Results (quality)



(a) Our method

(b) Reference (6 hours, 16 millions samples)

**Figure:** Result comparison for the Stanford Dragon model. Parameters for ketchup.

# Results (quality)



(a) Our result



(b) Reference (30 minutes,  $10^6$  samples)

**Figure:** Result comparison for the Happy Buddha model. Parameters for potato.

# Results (performance)

Model	# $\Delta$	Number of samples ( $N$ )			
		1	10	50	100
Bunny	$10^4$	2.1	5.3	19.8	38.2
Dragon	$10^5$	12.5	35.2	140.6	275.3
Buddha	$10^6$	96.7	97.7	128.0	216.0

**Table:** Timings in milliseconds of our method for different models and number of samples  $N$  (potato material properties). One directional light, 16 directions for rendering and reconstructing.

# Conclusions

What I did so far:

- Improved sampling (various corrections)
- Moved to a GPU random number generation (vs. texture-based)
- Introduced spectrally based computation

What are the next steps:

- Remove undesired color shifts
- Implement fast denoising for the final of the algorithm
- Comparison in performance with Optix ray-traced solution (special course)

# Conclusions



Thank you!

# References

- Eugene D'Eon and Geoffrey Irving. A quantized-diffusion model for rendering translucent materials. In *ACM SIGGRAPH 2011 Papers*, SIGGRAPH '11, pages 56:1--56:14, New York, NY, USA, 2011. ACM.
- J. R. Frisvad, T. Hachisuka, and T. K. Kjeldsen. Directional dipole for subsurface scattering in translucent materials. *ACM Transactions on Graphics*, 2014, -:--, 2014. To appear.
- Henrik Wann Jensen, Stephen R. Marschner, Marc Levoy, and Pat Hanrahan. A practical model for subsurface light transport. In *Proceedings of the 28th Annual Conference on Computer Graphics and Interactive Techniques*, SIGGRAPH '01, pages 511--518, New York, NY, USA, 2001. ACM.